

Quiz week 14

Question 5 (0 / 4 pt; mean 0.16)

Consider the following C code:

```
struct Files all_files[NUM_FILES];
int GetLastByteOfFile(int index) {
    if (index >= 0 && index < all_files) {
        struct File *file = &all_files[index];
        if (file->type == MEMORY) {
            return file->data[file->size - 1];
        } else if (file->type == DISK) {
            return GetLastByteOfDiskFile(file)
        } else {
            return -1;
        }
    } else {
        return -1;
    }
}
```

If the above function runs in kernel mode, we might be able to use a Spectre-style attack where the cache evictions caused by memory access of `file->data[file.size - 1]` allows us learn about the value of an arbitrary memory location. To perform this attack, the attacker would prefer to choose an out-of-bounds index such that ____.

- A. 88% ☐ the address of `all_files[index].data[file.size - 1]` is the memory address whose value they want to learn about
- B. 2% ☐ the address of `all_files[index].type` is the memory address they want to learn about
- C. 4% ^T
(correct) ☒ the address of `all_files[index].size` is the memory address they want to learn about
- D. ☐ the value of `all_files[index].type` is DISK

- `file = &all_files[index]`
- So if we provide an out of bounds index, we can read arbitrary memory
- `file->data[]`, ie `all_files[index].data[]` is like array2
- `file->size`, ie `all_files[index].size`, is like array1

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

our template

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

actual code

- kind ~ file->size, ie all_files[index].size, is like array1
 - This is the address we want to learn about by observing its cache behavior
- other_table [] ~ file->data[], ie all_files[index].data[] is like array2
 - Not using the .foo here

Question 5 (0 / 4 pt; mean 0.16)

Consider the following C code:

```
struct Files all_files[NUM_FILES];
int GetLastByteOfFile(int index) {
    if (index >= 0 && index < all_files) {
        struct File *file = &all_files[index];
        if (file->type == MEMORY) {
            return file->data[file->size - 1];
        } else if (file->type == DISK) {
            return GetLastByteOfDiskFile(file)
        } else {
            return -1;
        }
    } else {
        return -1;
    }
}
```

If the above function runs in kernel mode, we might be able to use a Spectre-style attack where the cache evictions caused by memory access of `file->data[file.size - 1]` allows us learn about the value of an arbitrary memory location. To perform this attack, the attacker would prefer to choose an out-of-bounds index such that ____.

- A. 88% ☐ the address of `all_files[index].data[file.size - 1]` is the memory address whose value they want to learn about
- B. 2% ☐ the address of `all_files[index].type` is the memory address they want to learn about
- C. 4% ^T
(correct) ☒ the address of `all_files[index].size` is the memory address they want to learn about
- D. ☐ the value of `all_files[index].type` is DISK

- Why not A?
- `all_files[index].data` is like `array2`
 - Based on which cache set is affected by different index values, we learn what those index values are
 - So we need to set up the index (~array1) to refer to the memory location we're interested in – here `file.size`, ie `all_files[index].size`

Q4

Consider the following code:

```
unsigned char check_array[32768];
int mystery = /* unknown */;

int Check(int key) {
    return check_array[(key + mystery) % 32768];
}
```

Suppose that:

- (to simplify the problem) virtual memory is not use
- check_array is located at physical address 0x1400000
- the system has a 2-way 64KB (2 to 16 byte) data cache with 64-byte cache blocks.
- Check is compiled to perform exactly three memory accesses:
 - to read the global variable read mystery
 - to reads its return address from the stack
 - to read from check_array

Suppose we determine that calling Check evicts from the data cache sets as follows:

key value	evicts values from cache set indexes
0	15, 72, 435
32	15, 72, 435
48	15, 72, 436
128	15, 72, 437
8240	15, 52, 72

Based on this information what is a possible value for mystery?

- $(0 + \text{mystery}) // 64 \pm \text{multiple of } 512 \text{ (number of cache sets)} = 435$
- $(32 + \text{mystery}) // 64 \pm \text{multiple of } 512 \text{ (number of cache sets)} = 435$
- $(48 + \text{mystery}) // 64 \pm \text{multiple of } 512 \text{ (number of cache sets)} = 436$
- ...

There was a typo originally that we fixed.
Originally wrote %16384 instead of %32768,
and 4096 instead of 8240 for the last row

Suppose we determine that calling Check evicts from the data cache sets as follows:

key value	evicts values from cache set indexes
0	15, 72, 435
32	15, 72, 435
48	15, 72, 436
128	15, 72, 437
8240	15, 52, 72

Based on this information what is a possible value for mystery?

Answer:

Key: a value between 27856 and 27871 +/- any multiple of 32768

originally we erroneously wrote % 16384 instead of % 32768, and 4096 instead of 8240 for the last value

$(0 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$

$(32 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$

$(48 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 436$

...

let's choose the multiple of 512 to be 0 for simplicity, for now

$\text{mystery} // 64 = 435$ implies mystery in $[435 * 64 = 27840, 27840 + 63 = 27903]$

$(32 + \text{mystery}) // 64 = 435$ implies mystery in $[435 * 64 - 32 = 27808, 27808 + 63 = 27871]$

$(48 + \text{mystery}) // 64 = 436$ implies mystery in $[436 * 64 - 48 = 27856, 27856 + 63 = 27919]$

overlap here implies mystery in $[27856, 27871]$

the multiple of 512 offsets this by $512 * 64 = 32768$